
CMPSC 473

Operating Systems Design & Construction

CPU Virtualization (cont.)
Dynamic Memory Allocation

Instructor: Ruslan Nikolaev

The Pennsylvania State University

September 7, 2023 – Lecture 6

Interrupts

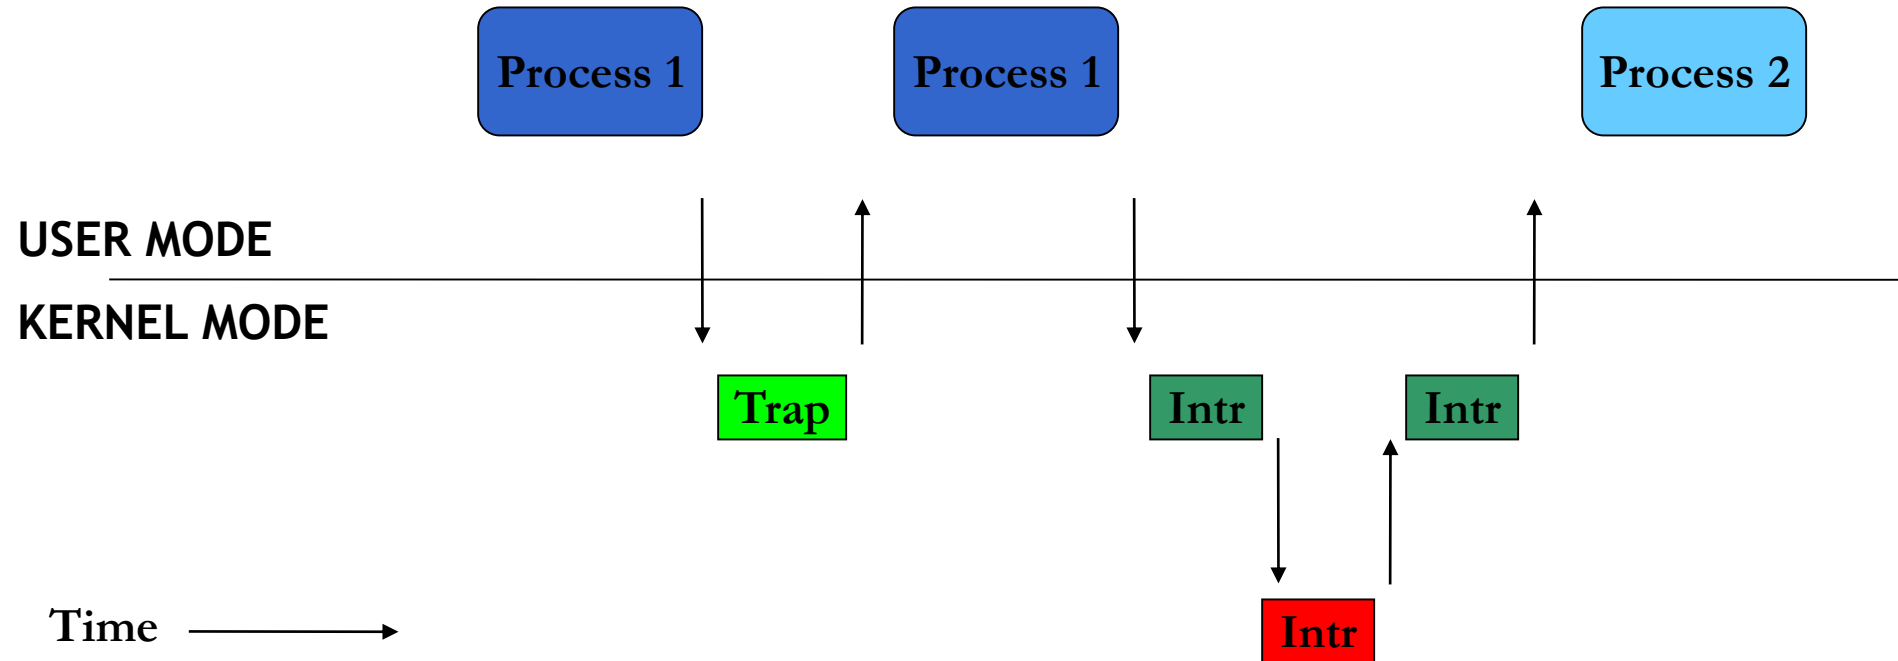
- CPU/OS design and operation for interrupts exactly like for traps
 - A CPU has its own well-defined set of interrupts
 - OS must implement a handler for each of these (part of OS code, just like trap handlers)
 - Table containing addresses of interrupt handles populated by OS during bootup, address of this table known to the CPU

Interrupts

OS @ boot (kernel mode)	Hardware	
initialize trap table	remember addresses of... syscall handler timer handler	
start interrupt timer	start timer interrupt CPU in X ms	
OS @ run (kernel mode)	Hardware	Program (user mode)
		Process A
		...
	timer interrupt save regs(A) \rightarrow k-stack(A) move to kernel mode jump to trap handler	
Handle the trap Call <code>switch()</code> routine save regs(A) \rightarrow proc_t(A) restore regs(B) \leftarrow proc_t(B) switch to k-stack(B) return-from-trap (into B)		
	restore regs(B) \leftarrow k-stack(B) move to user mode jump to B's PC	
		Process B
		...

Figure 6.3: Limited Direct Execution Protocol (Timer Interrupt)

Interrupts and Traps



- Only two ways to enter kernel mode from user mode
 - We previously already considered traps and trap handlers
 - Interrupts and corresponding interrupt handlers – the other way to enter kernel mode

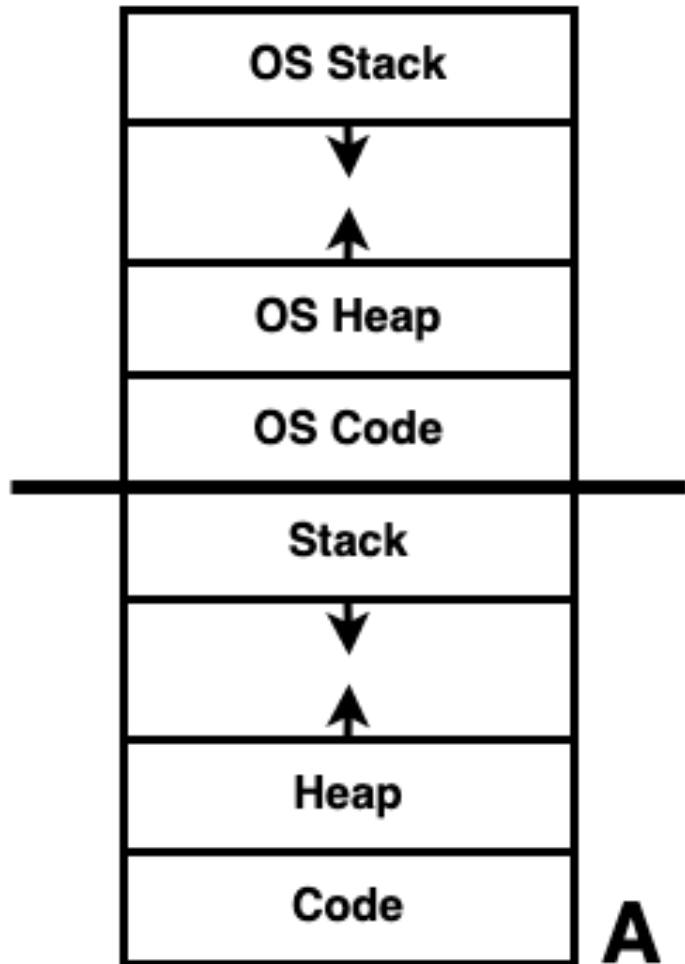
Question

Using recursion is generally considered undesirable in the design of trap/interrupt handlers. Why?

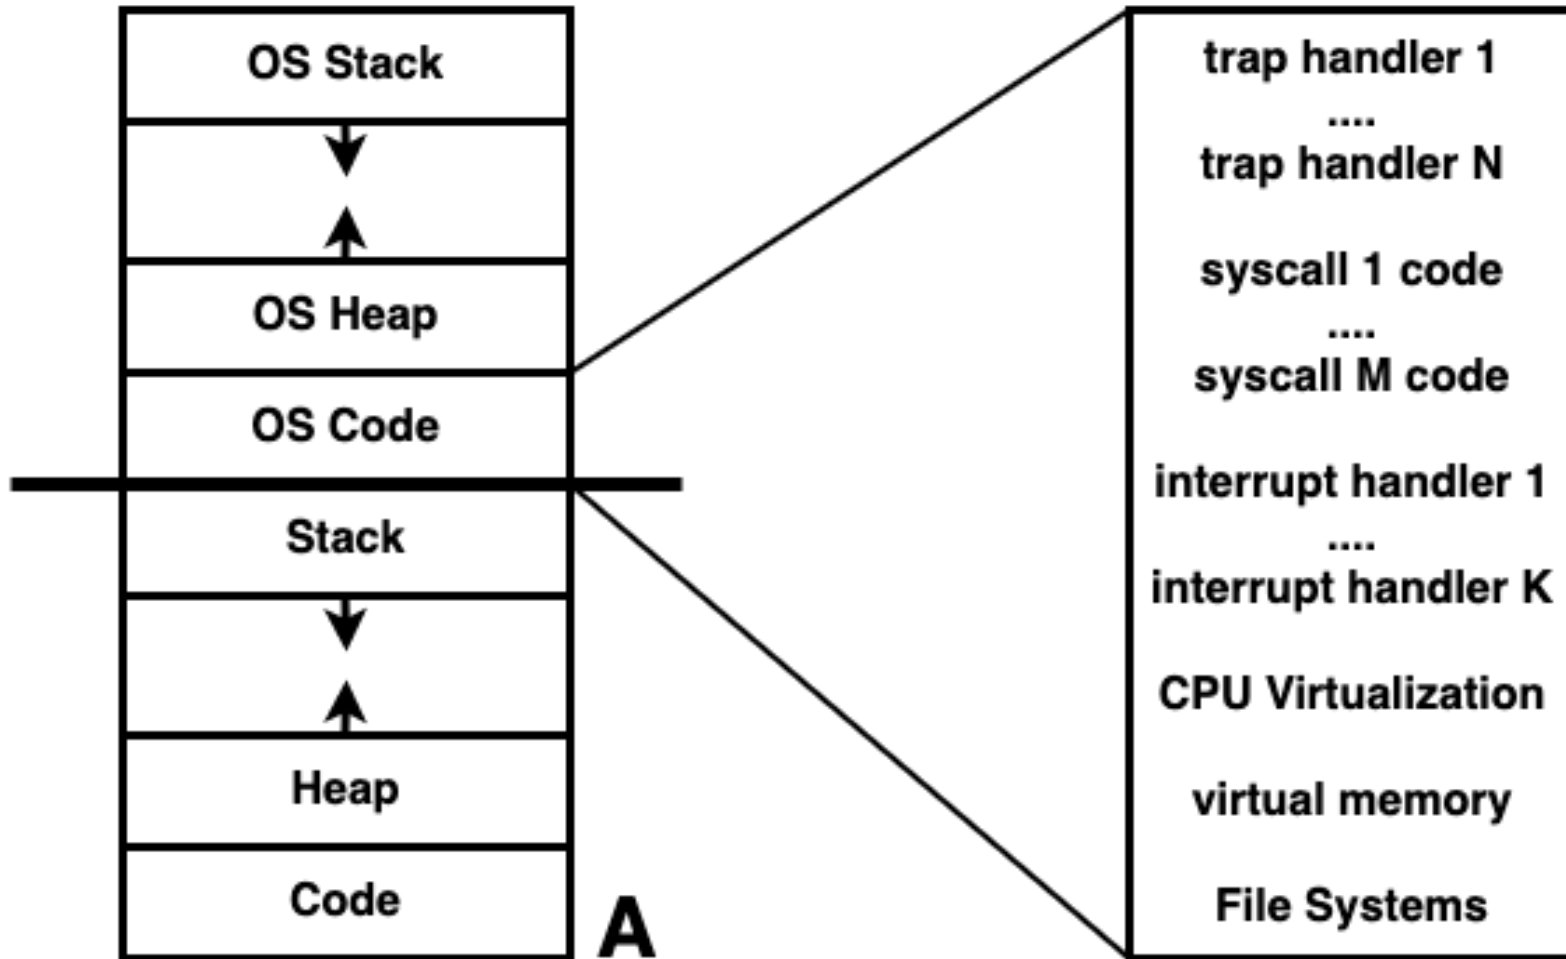
Signals vs. Interrupts/Traps

- Signals are application-level “interrupts”
- Creation: Two ways:
 - OS creates a signal for a process
 - A process creates a signal for another one with the OS help
 - kill system call
- Conveyance: by the OS
- Handling: Just like the OS implements interrupt handlers, a process implements signal handlers for all the signals defined by the OS


Signals vs. Interrupts/Traps



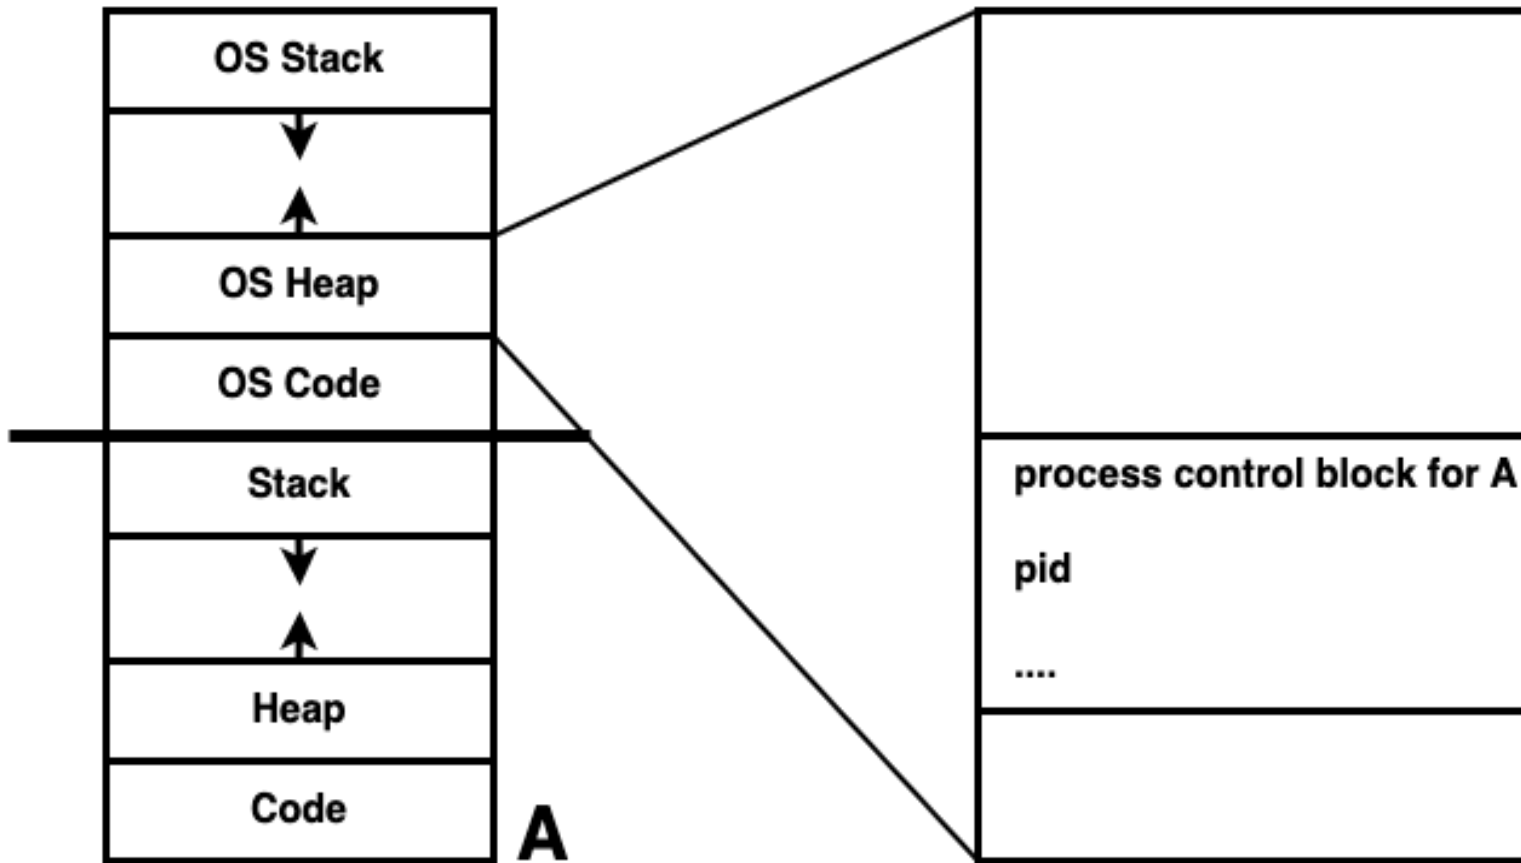
Signals vs. Interrupts/Traps



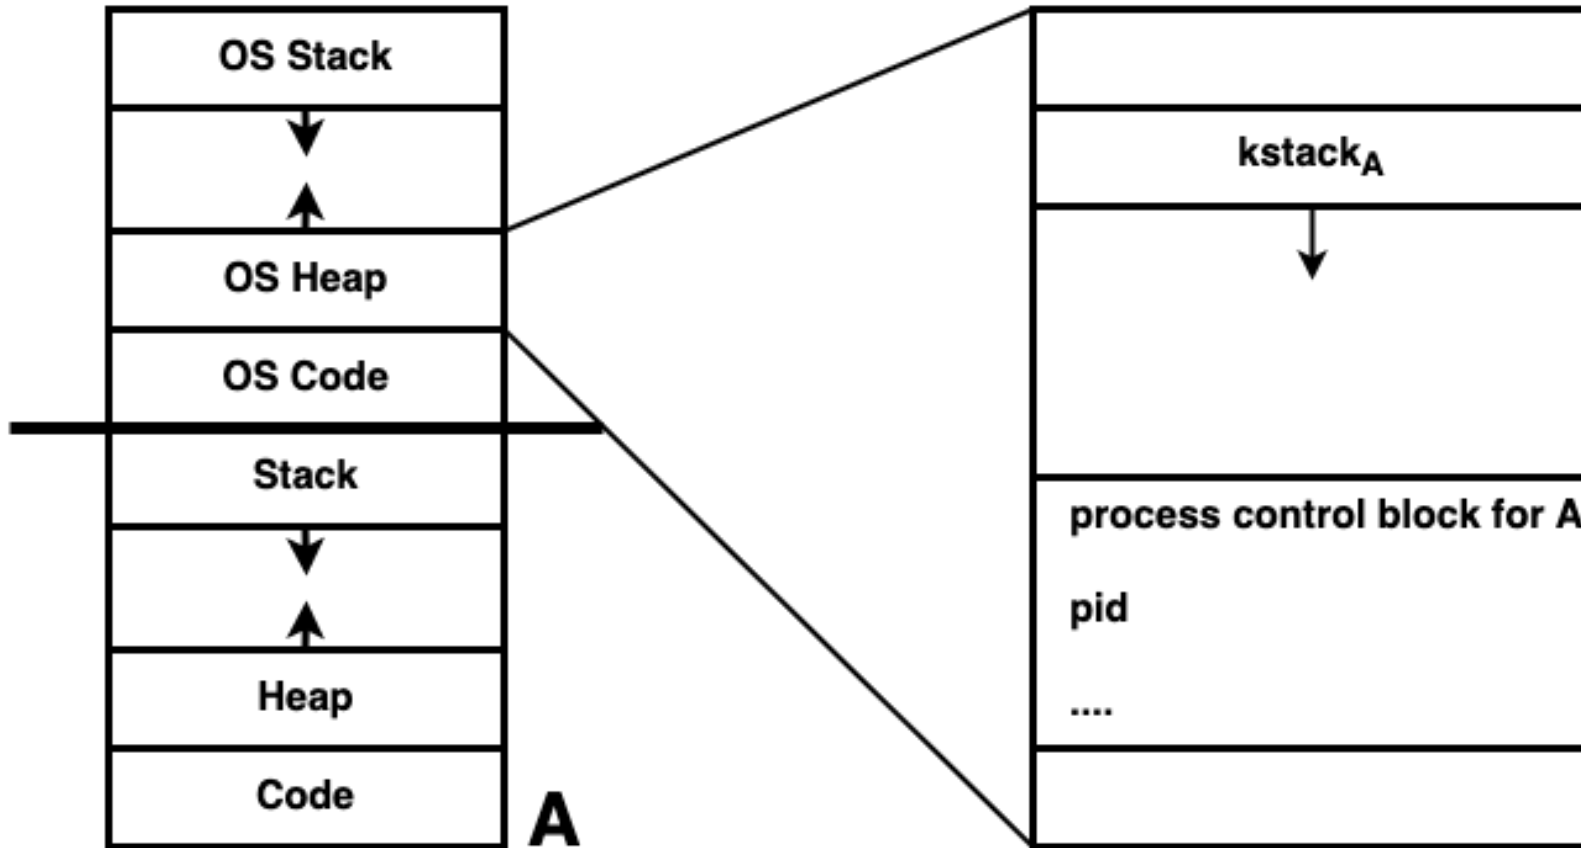
Signals vs. Interrupts/Traps



Signals vs. Interrupts/Traps



Signals vs. Interrupts/Traps

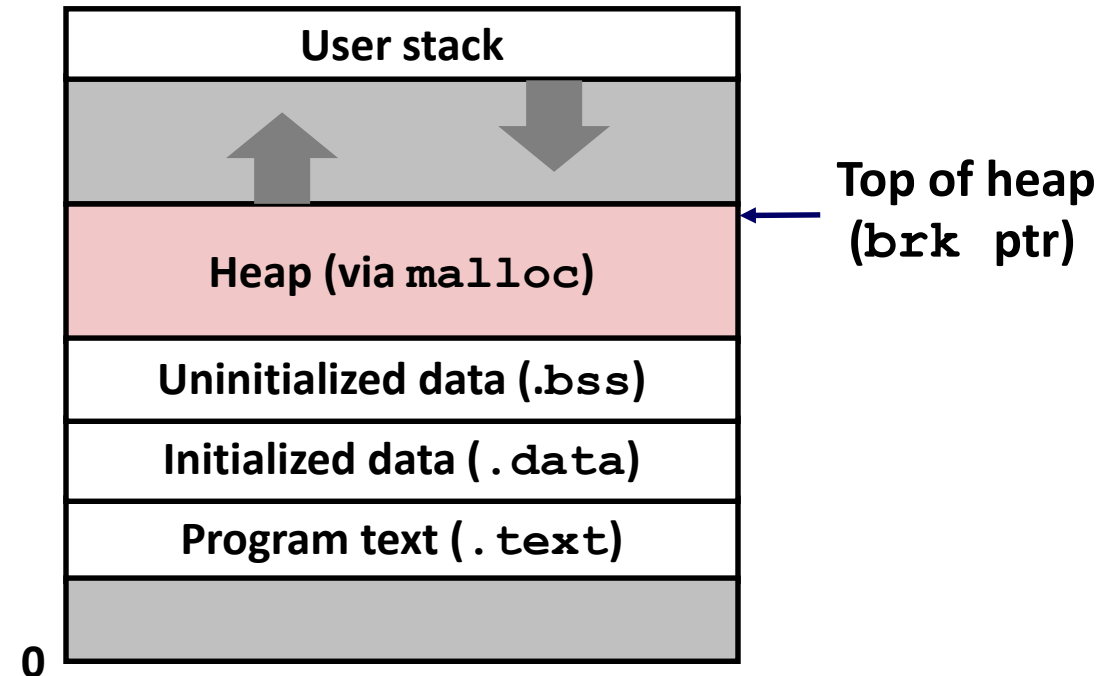



Overview

- Memory management occurs at 2 different granularities
 - OS virtual memory management
 - Virtual memory is procured from the OS at a relatively coarse granularity (page)
 - User-space dynamic memory allocation
 - Heap and stack management logic makes use of procured pages at a finer granularity (word)
 - We will study this first for the heap
-

Dynamic Memory Allocation

- Programmers use *dynamic memory allocators* (such as `malloc`) to acquire VM at run time.
 - For data structures whose size is only known at runtime.
- Dynamic memory allocators manage the *heap*.



Dynamic Memory Allocation

- Allocator maintains heap as collection of variable sized *blocks*, which are either *allocated* or *free*
- Types of allocators
 - ***Explicit allocator***: application allocates and frees space
 - E.g., `malloc` and `free` in C
 - ***Implicit allocator***: application allocates, but does not free space
 - E.g. garbage collection in Java, ML, and Lisp

The malloc Package

```
#include <stdlib.h>
```

```
void *malloc(size_t size) // see C notes on next slide
```

- **Successful:**
 - Returns a pointer to a memory block of at least **size** bytes aligned to an 8-byte (x86) or 16-byte (x86-64) boundary
 - If **size == 0**, returns NULL
- **Unsuccessful:** returns NULL (0) and sets **errno**

```
void free(void *p)
```

- Returns the block pointed at by **p** to pool of available memory
- **p** must come from a previous call to **malloc** or **realloc**

Other functions

- **calloc:** Version of **malloc** that initializes allocated block to zero.
 - **realloc:** Changes the size of a previously allocated block.
 - **sbrk:** Used internally by allocators to grow or shrink the heap
-

The malloc Package

```
#include <stdlib.h>
```

```
void *malloc(size_t size) // see C notes on next slide
```

- **Successful:**
 - Returns a pointer to a memory block of at least **size** bytes aligned to an 8-byte (x86) or 16-byte (x86-64) boundary
 - If **size == 0**, returns NULL
- **Unsuccessful:** returns NULL (0) and sets **errno**

```
void free(void *p)
```

- Returns the block pointed at by **p** to pool of available memory
- **p** must come from a previous call to **malloc** or **realloc**

Other functions

- **calloc:** Version of **malloc** that initializes allocated block to zero.
 - **realloc:** Changes the size of a previously allocated block.
 - **sbrk:** Used internally by allocators to grow or shrink the heap
-

Some C notes

- Void pointer
 - What?
 - No associated data type
 - Cannot dereference
 - Need to typecast before dereferencing
 - Why?
 - Need to return a dynamic data type, e.g., malloc
 - Need to implement an opaque object
 - Need to implement a generic function that will take different types of arguments determined at run-time

```
#include <stdlib.h>
```

```
void qsort(void *base, size_t nmemb, size_t size,  
           int (*compar)(const void *, const void *));
```

Some C notes

- Void pointer
 - What?
 - No associated data type
 - Cannot dereference
 - Need to typecast before dereferencing
 - Why?
 - Need to return a dynamic data type, e.g., malloc
 - Need to implement an opaque object
 - Need to implement a generic function that will take different types of arguments determined at run-time

```
#include <stdlib.h>
```

```
void qsort(void *base, size_t nmemb, size_t size,  
           int (*compar)(const void *, const void *));
```

Some C notes

- What is `size_t`?
- What is `errno`?
 - IMP: thread safe

malloc Example

```
#include <stdio.h>
#include <stdlib.h>

void foo(int n) {
    int i, *p;


    /* Allocate a block of n ints */
    p = (int *) malloc(n * sizeof(int));
    if (p == NULL) {
        perror("malloc");
        exit(0);
    }

    /* Initialize allocated block */
    for (i=0; i<n; i++)
        p[i] = i;

    /* Return allocated block to the heap */
    free(p);
}
```


Presentation Assumptions

- Memory is word addressed
- Each square is a word (e.g., 4 bytes for 32-bit x86)
 - Exceptions: in some slides, a square can be just 1 byte (for simplicity) when bytes are implied




Allocation Example


`p1 = malloc(4)`



`p2 = malloc(5)`




`p3 = malloc(6)`



`free(p2)`



`p4 = malloc(2)`



Each square is 1 byte here

Constraints

- Applications
 - Can issue arbitrary sequence of `malloc` and `free` requests
 - `free` request must be to a `malloc`'d block
 - Allocators
 - Cannot control number or size of allocated blocks
 - Must respond immediately to `malloc` requests
 - *i.e.*, cannot reorder or buffer requests
 - Must allocate blocks from free memory
 - *i.e.*, can only place allocated blocks in free memory
 - Can manipulate and modify only free memory
 - Must align blocks so they satisfy all alignment requirements
 - 8-byte (x86) or 16-byte (x86-64) alignment on Linux boxes
 - Cannot move the allocated blocks once they are `malloc`'d
 - *i.e.*, compaction is not allowed
-