

Operating Systems

CMPSC 473

Dynamic Memory Allocation

CPU/memory virtualization


September 19 2023 – Lecture 9

Instructor: Ruslan Nikolaev

ACK: some slides adapted from CSAPP text

Segregated List (Seglist) Allocators

- Each *size class* of blocks has its own free list



- Often have separate classes for each small size
- For larger sizes: One class for each two-power size

Seglist Allocator

- Given an array of free lists, each one for some size class
- To allocate a block of size n :
 - Search appropriate free list for block of size $m \geq n$
 - If an appropriate block is found:
 - Split block and place fragment on appropriate list (optional)
 - If no block is found, try next larger class
 - Repeat until block is found
- If no block is found:
 - Request additional heap memory from OS (using **sbrk()**)
 - Allocate block of n bytes from this new memory
 - Place remainder as a single free block in the corresponding size class

Seglist Allocator (cont.)

■ To free a block:

- Coalesce and place on appropriate list

■ Advantages of seglist allocators

- Higher throughput
 - log time for power-of-two size classes
- Better memory utilization
 - First-fit search of segregated free list approximates a best-fit search of entire heap.
 - Extreme case: Giving each block its own size class is equivalent to best-fit.

Memory-Related Perils and Pitfalls

- Dereferencing bad pointers
- Reading uninitialized memory
- Overwriting memory
- Referencing nonexistent variables
- Freeing blocks multiple times
- Referencing freed blocks
- Failing to free blocks

C operators

Operators

```
() [] -> .  
! ~ ++ -- + - * & (type) sizeof  
* / %  
+ -  
<< >>  
< <= > >=  
== !=  
&  
^  
|  
&&  
||  
?:  
= += -= *= /= %= &= ^= != <<= >>=  
,
```

Associativity

```
left to right  
right to left  
left to right  
left to right  
left to right  
left to right  
left to right  
left to right  
left to right  
right to left  
right to left  
left to right
```

- `->`, `()`, and `[]` have high precedence, with `*` and `&` just below
- Unary `+`, `-`, and `*` have higher precedence than binary forms

C Pointer Declarations: Test Yourself!

```
int *p
```

p is a pointer to int

```
int *p[13]
```

p is an array[13] of pointer to int

```
int *(p[13])
```

p is an array[13] of pointer to int

```
int **p
```

p is a pointer to a pointer to an int

```
int (*p)[13]
```

p is a pointer to an array[13] of int

```
int *f()
```

f is a function returning a pointer to int

```
int (*f)()
```

f is a pointer to a function returning int

```
int (*( *f()) [13])()
```

f is a function returning ptr to an array[13]
of pointers to functions returning int

```
int (*( *x[3]) ()) [5]
```

x is an array[3] of pointers to functions
returning pointers to array[5] of ints

Dereferencing Bad Pointers

■ The classic scanf bug

```
int val;  
  
...  
  
scanf("%d", val);
```


Reading Uninitialized Memory

- Assuming that heap data is initialized to zero

```
/* return y = Ax */
int *matvec(int **A, int *x) {
    int *y = malloc(N*sizeof(int));
    int i, j;

    for (i=0; i<N; i++)
        for (j=0; j<N; j++)
            y[i] += A[i][j]*x[j];
    return y;
}
```

To make sure the code works correctly without relying on the mentioned assumption, you should either initialize the y vector to zeros manually after allocating it:

```
int *y = malloc(N * sizeof(int));
for (int k = 0; k < N; k++) {
    y[k] = 0;
}
```

Overwriting Memory

- Allocating the (possibly) wrong sized object

```
int **p;  
  
p = malloc(N*sizeof(int)) ;  
  
for (i=0; i<N; i++) {  
    p[i] = malloc(M*sizeof(int)) ;  
}
```

```
int **p;  
p = malloc(N * sizeof(int*));  
for (i=0; i<N; i++) {  
    p[i] = malloc(M*sizeof(int));  
}
```

Overwriting Memory

■ Off-by-one error

```
int **p;  
  
p = malloc(N*sizeof(int *));  
  
for (i=0; i<=N; i++) {  
    p[i] = malloc(M*sizeof(int));  
}
```

```
for (i=0; i<N; i++) {  
    p[i] = malloc(M*sizeof(int));  
}
```

Overwriting Memory

- Not checking the max string size

```
char s[8];  
int i;  
  
gets(s);  /* reads "123456789" from stdin */
```

- Basis for classic buffer overflow attacks

Overwriting Memory

■ Misunderstanding pointer arithmetic

```
int *search(int *p, int val) {  
  
    while (*p && *p != val)  
        p += sizeof(int);  
  
    return p;  
}
```

```
int *search(int *p, int val) {  
    while (*p && *p != val)  
        p++;  
    return p;  
}
```

Overwriting Memory

- Referencing a pointer instead of the object it points to

```
int *BinheapDelete(int **binheap, int *size) {  
    int *packet;  
    packet = binheap[0];  
    binheap[0] = binheap[*size - 1];  
    *size--;  
    Heapify(binheap, *size, 0);  
    return(packet);  
}
```

```
int *BinheapDelete(int **binheap, int *size) {  
    int *packet;  
    packet = binheap[0];  
    binheap[0] = binheap[*size - 1];  
    (*size)--;  
    Heapify(binheap, *size, 0);  
    return packet;  
}
```

Referencing Nonexistent Variables

- Forgetting that local variables disappear when a function returns

```
int *foo () {  
    int val;  
  
    return &val;  
}
```

Freeing Blocks Multiple Times

■ Nasty!

```
x = malloc(N*sizeof(int));  
    <manipulate x>  
free(x);  
  
y = malloc(M*sizeof(int));  
    <manipulate y>  
free(x);
```


Referencing Freed Blocks

■ Evil!

```
x = malloc(N*sizeof(int));  
  <manipulate x>  
free(x);  
  ...  
y = malloc(M*sizeof(int));  
for (i=0; i<M; i++)  
    y[i] = x[i]++;
```

Failing to Free Blocks (Memory Leaks)

■ Slow, long-term killer!

```
foo() {  
    int *x = malloc(N*sizeof(int));  
    ...  
    return;  
}
```

The pointer `x` is allocated memory dynamically but is never freed before the function `foo()` returns. When the function returns, the local pointer `x` goes out of scope, and we lose the reference to the allocated memory, leading to a memory leak.

To fix the memory leak, you need to free the allocated memory before the function returns:

```
foo() {  
    int *x = malloc(N*sizeof(int));  
    ...  
    free(x);  
    return;  
}
```

This loop iteratively frees each node in the list by temporarily holding a reference to the current node, moving the head pointer to the next node, and then freeing the temporary node.

The corrected foo() function would be:

Failing to Free Blocks (Memory Leaks)

■ Freeing only part of a data structure

```
struct list {  
    int val;  
    struct list *next;  
};
```

```
foo() {  
    struct list *head = malloc(sizeof(struct list));  
    head->val = 0;  
    head->next = NULL;  
    <create and manipulate the rest of the list>  
    ...  
    free(head);  
    return;  
}
```

```
C  
// copy code  
foo() {  
    struct list *head = malloc(sizeof(struct list));  
    head->val = 0;  
    head->next = NULL;  
    // <create and manipulate the rest of the list>  
  
    while (head != NULL) {  
        struct list *temp = head;  
        head = head->next;  
        free(temp);  
    }  
    return;  
}
```