


# Operating Systems

## CMPSC 473

### CPU Virtualization - Scheduling


#### Lecture 14: October 5, 2023

Instructor: Ruslan Nikolaev



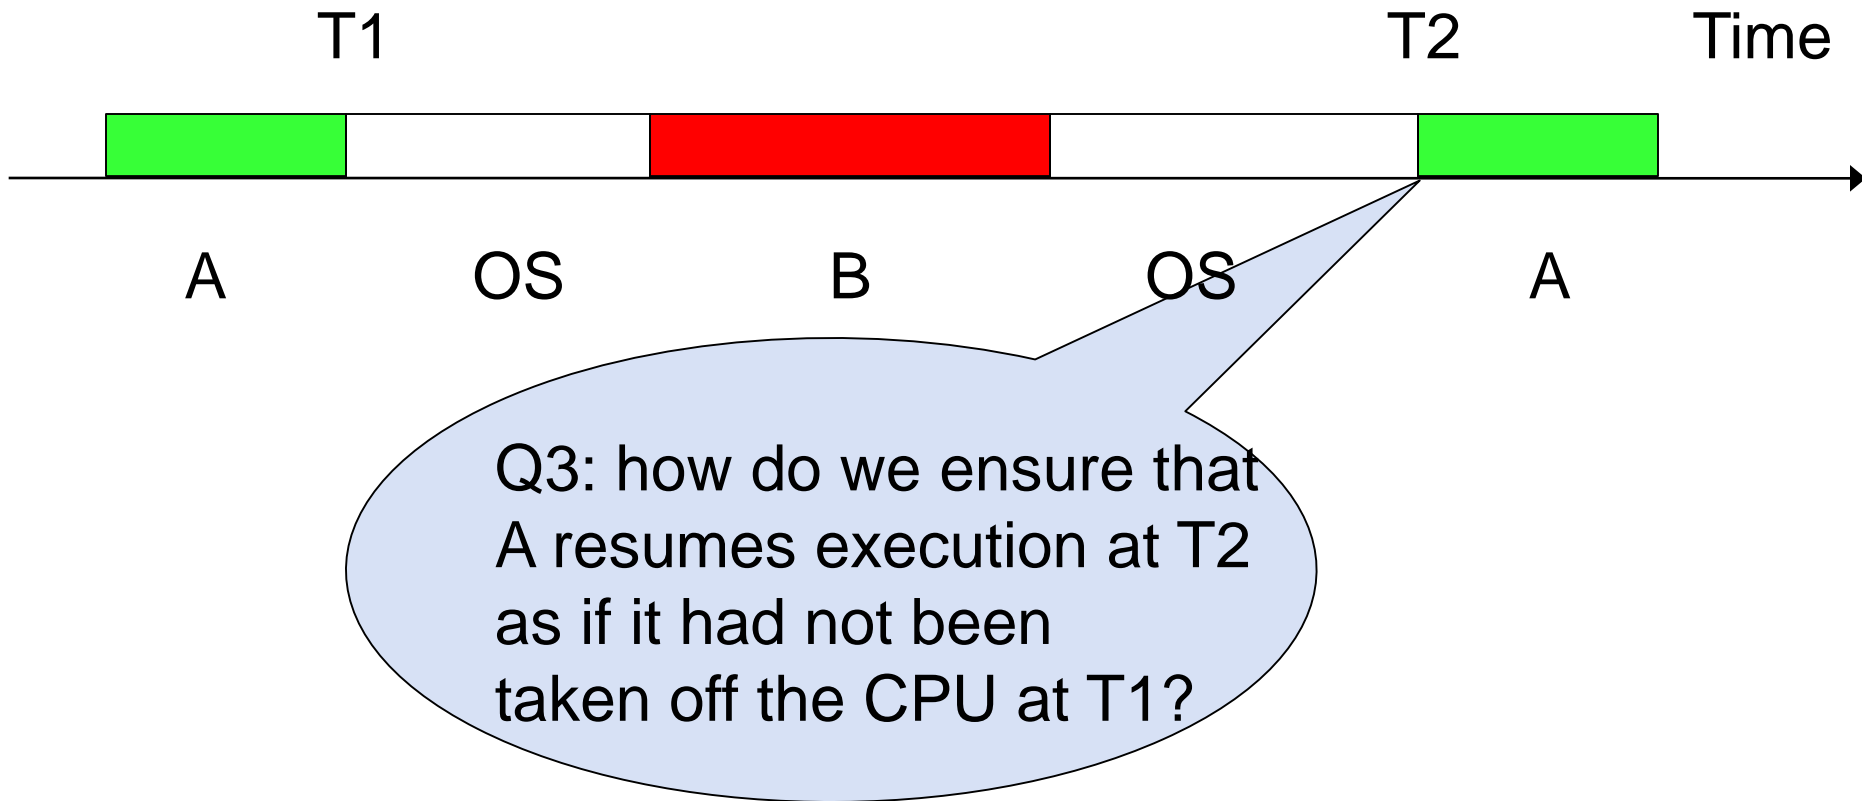
Q1: what if the process does something undesirable here?

- OS support
  - Trap handlers, signal handling
  - Per-process kernel stack
  - Virtual memory, page tables
- Hardware support
  - Trap mechanism, CPU modes
  - Switching to the specified (per-CPU) kernel stack, saving a user stack register, and program counter
  - MMU, page tables, TLB




Q2: how does the OS get to start running here?

- OS support
  - Interrupt handlers
  - Per-process kernel stack
- Hardware support
  - Interrupt mechanism, e.g., a timer interrupt
  - Switching to the specified (per-CPU) kernel stack, saving a user stack register, and program counter



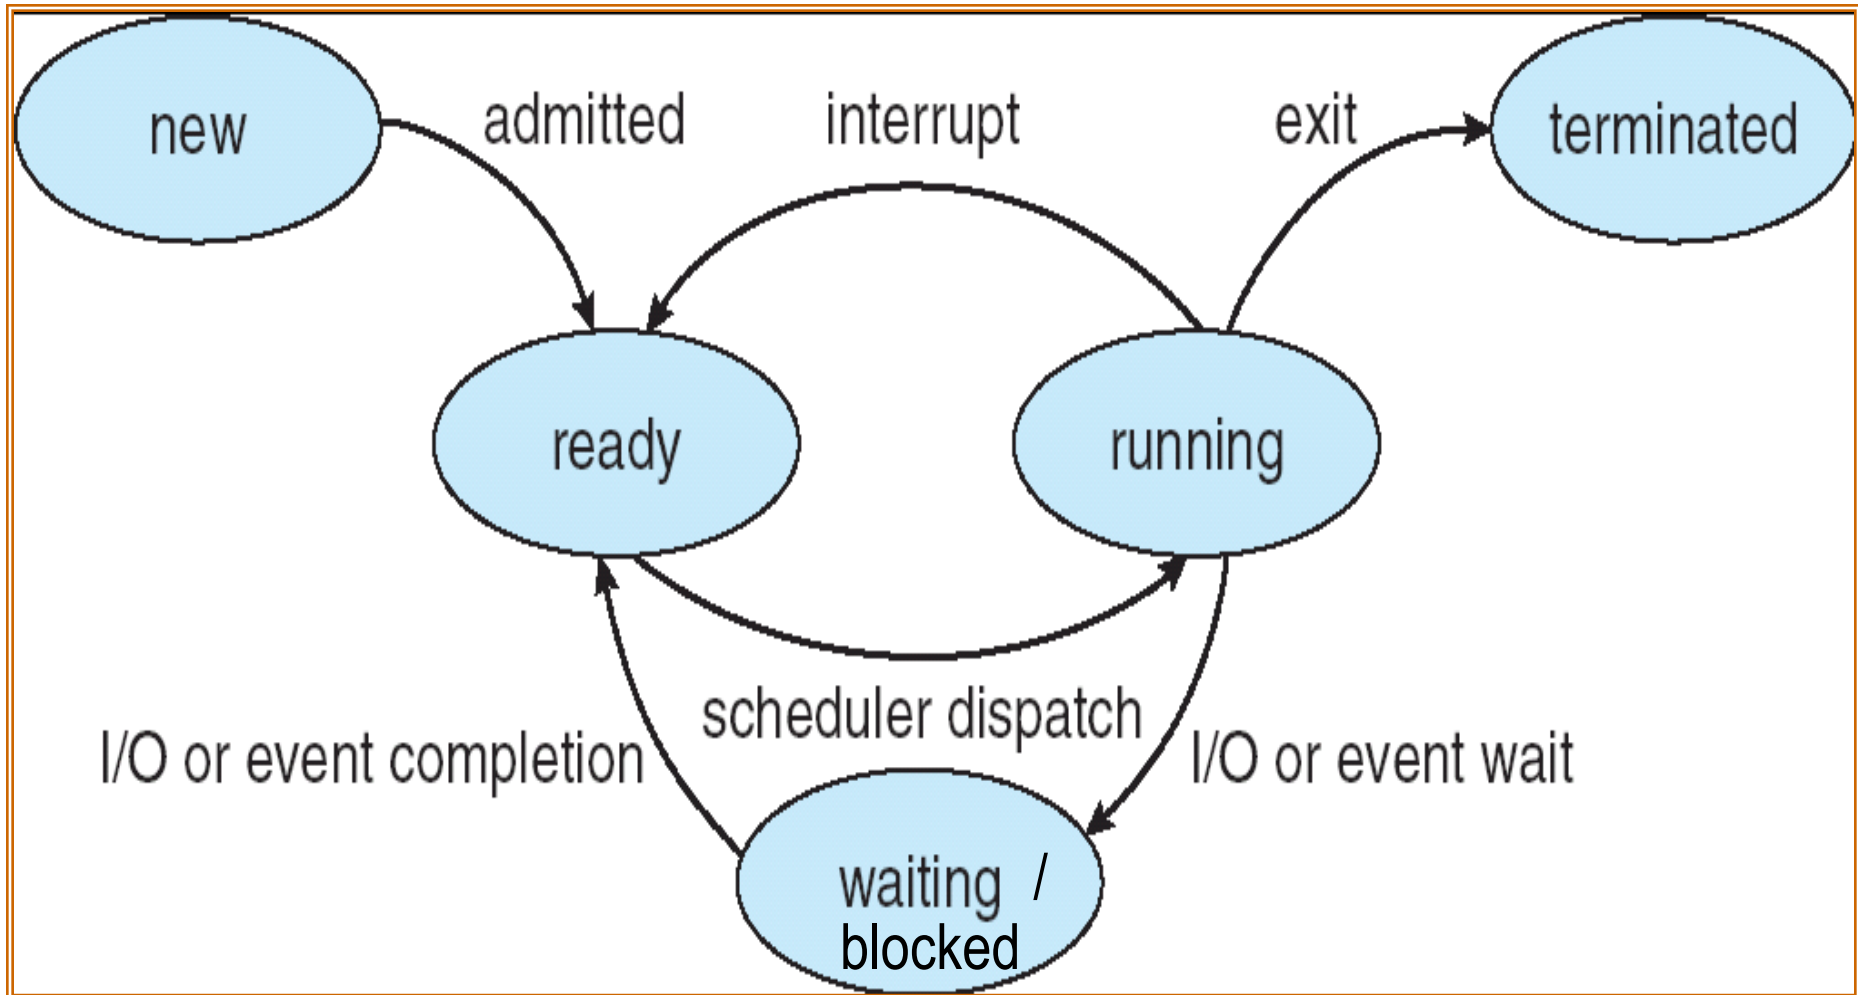
- OS support
  - Saving/restoring registers in/from PCB (process control block)
- Hardware support
  - Not necessary unless using “hardware context switching”


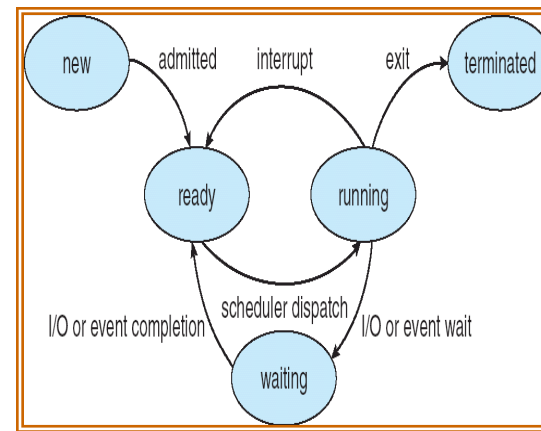


Q4: How does the OS decide which process to run next?


- The CPU scheduler

# Scheduling states of a process/thread







Running




Ready


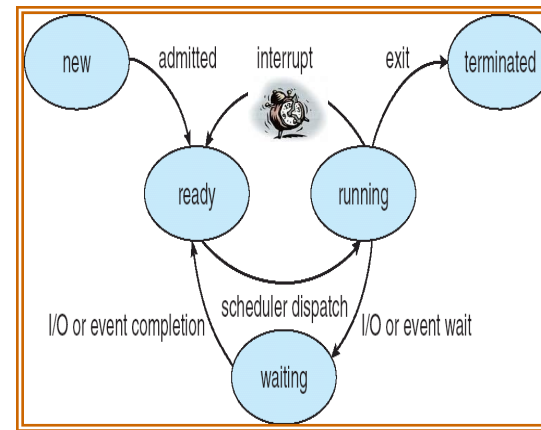
Lock




Waiting





Timer interrupt




Lock









Running




Ready


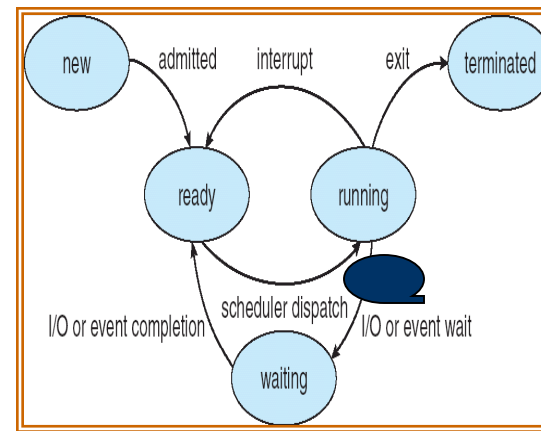



Waiting


Lock 





I/O system call



Running




Ready





Waiting

Lock



Disk

Lets pick the second  
process in the ready  
queue



# CPU scheduler: important concerns

- Optimization metric/criterion
  - Latency metrics:
    - Turnaround/completion time: time between when a “job” is submitted and when it finishes
    - Response time: time between when a process desires the CPU (“ready”) and when it gets to run on the CPU (“running”)
  - Throughput metrics: amount of “work” per unit time
    - e.g., # of requests or processes finishing per unit time
  - “Fairness”
    - Proportional-share
      - A job has a weight and gets resources proportional to its weight
    - Max-min
    - Many others
  - Combinations of these
- Overheads of the algorithm itself
  - Runtime and space complexity

# FIFO/FCFS

- Overheads
  - Picking the next process to run
  - Process arrival, process departure
- How does it do for these metrics?
  - Response time
  - Throughput
  - Fairness
- Pros:
  - Simple to implement! Low overheads
  - Does well for “batch” jobs
- Cons:
  - Short jobs will suffer, a large job will monopolize the CPU

# Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time
- SJF is optimal for avg. waiting time – gives minimum average waiting time for a given set of processes

# Why Preemption is Necessary

- To improve CPU utilization
  - Most processes are not *ready* at all times during their lifetimes
  - E.g., think of a text editor waiting for input from the keyboard
  - Also improves I/O utilization
- To improve responsiveness
  - Many processes would prefer to receive CPU quickly when they need it
- Modern CPU schedulers are preemptive

## SJF: Variations on the theme

- **Non-preemptive:** once CPU given to the process it cannot be preempted until completes its CPU burst - the SJF we already saw
- **Preemptive:** if a new process arrives with CPU length less than remaining time of current executing process, preempt
  - This scheme is known as *Shortest-Remaining-Time-First (SRTF)*
  - Also called *Shortest Remaining Processing Time (SRPT)*
- **Overheads?**
  - Compare using unordered linked list vs. ordered list vs. heap
- **Why SJF/SRTF may not be practical**
  - CPU requirement of a process rarely known in advance



# Round Robin (RR)

- Each process gets a small unit of CPU time (*time quantum*), usually 1-10 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue
- If there are  $n$  processes in the ready queue and the time quantum is  $q$ , then each process gets  $1/n$  of the CPU time in chunks of at most  $q$  time units at once. No process waits more than  $(n-1)q$  time units
- Performance
  - $q$  large  $\Rightarrow$  FCFS
  - $q$  small  $\Rightarrow q$  must be large with respect to the context switch cost, otherwise overhead is too high


# Proportional-Share Schedulers

- A generalization of round robin
- Process  $P_i$  given a CPU weight  $w_i > 0$

# Lottery Scheduling

- Perhaps the simplest proportional-share scheduler
- Create lottery tickets equal to the sum of the weights of all processes
- Draw a lottery ticket and schedule the process that owns that ticket

# Lottery Scheduling Example



9

$P1=6$


|   |   |
|---|---|
| 1 | 4 |
| 2 | 5 |
| 3 | 6 |

$P2=9$

|   |    |    |
|---|----|----|
| 7 | 10 | 13 |
| 8 | 11 | 14 |
| 9 | 12 | 15 |

*Schedule P2*

# Lottery Scheduling Example



3

$P1=6$


|   |   |
|---|---|
| 1 | 4 |
| 2 | 5 |
| 3 | 6 |

$P2=9$

|   |    |    |
|---|----|----|
| 7 | 10 | 13 |
| 8 | 11 | 14 |
| 9 | 12 | 15 |

*Schedule P1*

# Lottery Scheduling Example



11

$P1=6$

|   |   |
|---|---|
| 1 | 4 |
| 2 | 5 |
| 3 | 6 |

$P2=9$

|   |    |    |
|---|----|----|
| 7 | 10 | 13 |
| 8 | 11 | 14 |
| 9 | 12 | 15 |

- As  $t \rightarrow \infty$ , processes will get their share (unless they were blocked a lot)
- Problem with Lottery scheduling: Only probabilistic guarantee
- What does the scheduler have to do
  - When a new process arrives?
  - When a process terminates?

*Schedule P2*

## Multi-Level Feedback Queue (MLFQ)

- See Chapter 8 of OSTEP
- Optional reading, not on syllabus
- Very old idea, a great example of heuristics that
  - Try to combine conflicting goals of good response times for interactive jobs and fairness
  - Learn dynamically the nature of a job (whether it is interactive)